

AB_Test

January 17, 2025

```
[56]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats as stats

from scipy.stats import chi2_contingency
from scipy.stats import mannwhitneyu

from sklearn.impute import KNNImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.metrics import classification_report, accuracy_score, roc_auc_score
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, cross_val_score

from statsmodels.stats.outliers_influence import variance_inflation_factor
```

1 Introduction

In this project I will be analyzing a bank loan marketing conversion dataset obtained from Kaggle (<https://www.kaggle.com/datasets/arashnic/banking-loan-prediction/data>). The dataset features attributes that pertain to a bank loan application. The dataset includes several columns regarding the applicants background, the source of the application, terms of the bank loan, and whether or not the loan was approved.

I will be conducting several A/B tests to identify the most effective lead generating source category with regards to approval rate and loan amount. I will create and test a predictive model using logistic regression.

1.0.1 Exploratory Data Analysis

Dataset Overview:

```
[57]: df = pd.read_csv("../data/train_loan/train.csv")

print(f"This dataset has {df.shape[0]} rows and {df.shape[1]} attributes. ")

print(df.describe())
```

```
df.head()
```

This dataset has 69713 rows and 22 attributes.

	Employer_Category2	Monthly_Income	Existing_EMI	Loan_Amount	\
count	65415.000000	6.971300e+04	69662.000000	42004.000000	
mean	3.720187	5.622283e+03	360.928751	39429.982859	
std	0.807374	1.747671e+05	2288.517927	30727.595990	
min	1.000000	0.000000e+00	0.000000	5000.000000	
25%	4.000000	1.650000e+03	0.000000	20000.000000	
50%	4.000000	2.500000e+03	0.000000	30000.000000	
75%	4.000000	4.000000e+03	350.000000	50000.000000	
max	4.000000	3.838384e+07	545436.500000	300000.000000	

	Loan_Period	Interest_Rate	EMI	Var1	Approved
count	42004.000000	22276.000000	22276.000000	69713.000000	69713.000000
mean	3.890629	19.213570	1101.466242	3.948446	0.014631
std	1.167491	5.847136	752.661394	3.819214	0.120073
min	1.000000	11.990000	118.000000	0.000000	0.000000
25%	3.000000	15.250000	649.000000	0.000000	0.000000
50%	4.000000	18.000000	941.000000	2.000000	0.000000
75%	5.000000	20.000000	1295.000000	7.000000	0.000000
max	6.000000	37.000000	13556.000000	10.000000	1.000000

[57]:

	ID	Gender	DOB	Lead_Creation_Date	City_Code	\
0	APPC90493171225	Female	23/07/79	15/07/16	C10001	
1	APPD40611263344	Male	07/12/86	04/07/16	C10003	
2	APPE70289249423	Male	10/12/82	19/07/16	C10125	
3	APPF80273865537	Male	30/01/89	09/07/16	C10477	
4	APPG60994436641	Male	19/04/85	20/07/16	C10002	

	City_Category	Employer_Code	Employer_Category1	Employer_Category2	\
0	A	COM0044082	A	4.0	
1	A	COM0000002	C	1.0	
2	C	COM0005267	C	4.0	
3	C	COM0004143	A	4.0	
4	A	COM0001781	A	4.0	

	Monthly_Income	...	Contacted	Source	Source_Category	Existing_EMI	\
0	2000.0	...	N	S122	G	0.0	
1	3500.0	...	Y	S122	G	0.0	
2	2250.0	...	Y	S143	B	0.0	
3	3500.0	...	Y	S143	B	0.0	
4	10000.0	...	Y	S134	B	2500.0	

	Loan_Amount	Loan_Period	Interest_Rate	EMI	Var1	Approved
0	NaN	NaN	NaN	NaN	0	0
1	20000.0	2.0	13.25	953.0	10	0

2	45000.0	4.0	NaN	NaN	0	0
3	92000.0	5.0	NaN	NaN	7	0
4	50000.0	2.0	NaN	NaN	10	0

[5 rows x 22 columns]

Data Cleaning:

```
[58]: #Removing any duplicate rows
df.drop_duplicates(inplace=True)

print(df.info())

null_rows_sum = df.isnull().any(axis=1).sum()
print(f"There are {null_rows_sum} rows with atleast one null value in a column.
↪")
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 69713 entries, 0 to 69712

Data columns (total 22 columns):

#	Column	Non-Null Count	Dtype
0	ID	69713 non-null	object
1	Gender	69713 non-null	object
2	DOB	69698 non-null	object
3	Lead_Creation_Date	69713 non-null	object
4	City_Code	68899 non-null	object
5	City_Category	68899 non-null	object
6	Employer_Code	65695 non-null	object
7	Employer_Category1	65695 non-null	object
8	Employer_Category2	65415 non-null	float64
9	Monthly_Income	69713 non-null	float64
10	Customer_Existing_Primary_Bank_Code	60322 non-null	object
11	Primary_Bank_Type	60322 non-null	object
12	Contacted	69713 non-null	object
13	Source	69713 non-null	object
14	Source_Category	69713 non-null	object
15	Existing_EMI	69662 non-null	float64
16	Loan_Amount	42004 non-null	float64
17	Loan_Period	42004 non-null	float64
18	Interest_Rate	22276 non-null	float64
19	EMI	22276 non-null	float64
20	Var1	69713 non-null	int64
21	Approved	69713 non-null	int64

dtypes: float64(7), int64(2), object(13)

memory usage: 11.7+ MB

None

There are 49405 rows with atleast one null value in a column.

Many of the attributes have null values in 1 or more rows, however since the initial A/B tests will only make use of the Approved and Source_Category columns, which don't have any null values, we do not need to discard or impute any rows for now.

Next I will check for outliers, before doing so I must convert the DOB and Lead_Creation_Date columns to Age, and Lead_Age.

```
[59]: #Converting lead creation date and DOB to datetime
df["Lead_Creation_Date"] = pd.to_datetime(df["Lead_Creation_Date"])
df["DOB"] = pd.to_datetime(df["DOB"])

#Finding most recent loan creation date
most_recent_date = df["Lead_Creation_Date"].max()

#Calculating loan application age in years
df["Lead_Age"] = (most_recent_date - df["Lead_Creation_Date"]).dt.days / 365

#Calculating age of applicants in years
df["Age"] = (most_recent_date - df["DOB"]).dt.days / 365

print(df["Age"].describe())
print(df["Lead_Age"].describe())

max_dob = df["DOB"].max()
print(max_dob)
```

```
C:\Users\dimdi\AppData\Local\Temp\ipykernel_13716\1392067781.py:2: UserWarning:
Could not infer format, so each element will be parsed individually, falling
back to `dateutil`. To ensure parsing is consistent and as-expected, please
specify a format.
```

```
df["Lead_Creation_Date"] = pd.to_datetime(df["Lead_Creation_Date"])
```

```
C:\Users\dimdi\AppData\Local\Temp\ipykernel_13716\1392067781.py:3: UserWarning:
Could not infer format, so each element will be parsed individually, falling
back to `dateutil`. To ensure parsing is consistent and as-expected, please
specify a format.
```

```
df["DOB"] = pd.to_datetime(df["DOB"])
```

```
count    69698.000000
mean      22.841209
std       22.893110
min      -58.093151
25%       25.284932
50%       28.465753
75%       32.254795
max       41.967123
Name: Age, dtype: float64
count    69713.000000
mean      0.365079
std       0.207192
```

```

min            0.000000
25%            0.224658
50%            0.312329
75%            0.405479
max            0.923288
Name: Lead_Age, dtype: float64
2074-12-29 00:00:00

```

We can see that the new Age column has a minimum value of -58, after checking the maximum date of birth, we can see that one of the applicants has their birth year set to 2074, which is clearly an error.

I will now replace the outliers using the interquartile range method. To better preserve the scope of the data, I will raise the threshold value for the IQR range from a standard 1.5 to 4.5. This way only the most extreme outliers will be set to null and later imputed.

I will take a different approach with the age column however, instead of using the IQR range I will simply set values under 18 years to null, since the current maximum age of 42 years is acceptable.

```

[60]: numerical_columns = ["Monthly_Income", "Existing_EMI", "Loan_Amount",
    ↪ "Loan_Period", "Interest_Rate", "EMI", "Age", "Lead_Age"]

for col in df[numerical_columns]:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower_bound = Q1 - (4.5 * IQR)
    upper_bound = Q3 + (4.5 * IQR)

    if col == "Age":
        num_outliers = ((df[col] < 18)).sum()
        print(f"Column {col}: {num_outliers} outliers replaced with NaN.")
        df[col] = df[col].mask((df[col] < 18))
    else:
        num_outliers = ((df[col] < lower_bound) | (df[col] > upper_bound)).sum()
        print(f"Column {col}: {num_outliers} outliers replaced with NaN.")
        df[col] = df[col].mask((df[col] < lower_bound) | (df[col] >
    ↪ upper_bound))

df[numerical_columns].describe()

```

```

Column Monthly_Income: 1729 outliers replaced with NaN.
Column Existing_EMI: 3211 outliers replaced with NaN.
Column Loan_Amount: 124 outliers replaced with NaN.
Column Loan_Period: 0 outliers replaced with NaN.
Column Interest_Rate: 0 outliers replaced with NaN.
Column EMI: 154 outliers replaced with NaN.
Column Age: 5833 outliers replaced with NaN.

```

Column Lead_Age: 0 outliers replaced with NaN.

```
[60]:
```

	Monthly_Income	Existing_EMI	Loan_Amount	Loan_Period	\
count	67984.000000	66451.000000	41880.000000	42004.000000	
mean	3035.341438	196.711859	38884.765998	3.890629	
std	2122.256046	383.288838	29051.539138	1.167491	
min	0.000000	0.000000	5000.000000	1.000000	
25%	1600.000000	0.000000	20000.000000	3.000000	
50%	2500.000000	0.000000	30000.000000	4.000000	
75%	3800.000000	250.000000	50000.000000	5.000000	
max	14550.000000	1923.000000	185000.000000	6.000000	

	Interest_Rate	EMI	Age	Lead_Age
count	22276.000000	22122.000000	63865.000000	69713.000000
mean	19.213570	1073.257752	29.607387	0.365079
std	5.847136	666.681053	4.672571	0.207192
min	11.990000	118.000000	18.013699	0.000000
25%	15.250000	639.000000	26.131507	0.224658
50%	18.000000	934.000000	29.041096	0.312329
75%	20.000000	1276.000000	32.668493	0.405479
max	37.000000	4194.000000	41.967123	0.923288

Outliers have been removed, upon looking at the description of the data we can see that all min, max and mean values are acceptable.

Before proceeding to feature engineering and visualization, I will drop unneeded and inapplicable columns as they will slow down computation time and interfere with the predictive modeling. To build the predictive model using logistic regression, I will have to factor all categorical columns to into binary columns using One Hot Encoding. This algorithm will encode the categorical variables by creating a new binary column for each unique value for each categorical variable. Depending on the amount of unique values in the categorical variables, this method may massively overfit the logistic regression model. Therefore it would wise to drop categorical columns that have too many unique values. This will assist in making the visualization of the categorical columns easier to interpret.

```
[61]: categorical_columns = ["Gender", "City_Code", "City_Category", "Employer_Code",  
    ↪ "Employer_Category1", "Employer_Category2",  
    ↪ "Customer_Existing_Primary_Bank_Code", "Primary_Bank_Type", "Contacted",  
    ↪ "Source", "Source_Category"]  
  
for col in categorical_columns:  
    unique_count = df[col].nunique()  
    print(f" Unique Values in column {col} = {unique_count}")
```

```
Unique Values in column Gender = 2  
Unique Values in column City_Code = 678  
Unique Values in column City_Category = 3  
Unique Values in column Employer_Code = 36617  
Unique Values in column Employer_Category1 = 3
```

```

Unique Values in column Employer_Category2 = 4
Unique Values in column Customer_Existing_Primary_Bank_Code = 57
Unique Values in column Primary_Bank_Type = 2
Unique Values in column Contacted = 2
Unique Values in column Source = 29
Unique Values in column Source_Category = 7

```

We can see that City_Code, Employer_Code, Customer_Existing_Primary_Bank_Code and Source have the highest cardinalities. These columns will be removed, as introducing that many individual columns into the logistic regression will overfit the model. The columns to be removed also each have their own respective category columns (“City_Category”, “Employer_Category1”, “Employer_Category2”, “Primary_Bank_Type” and “Source_Category”) which will do a good job of representing the removed columns in the model.

I will also drop the ID column since it is simply the customers ID and does not help in determining approval. The “Var1” column will also be dropped, as it is not clear what it represents due to its ambiguous title and Kaggle description: “Anonymized Categorical variable with multiple levels”.

```

[62]: df = df.drop(columns=["ID", "Var1", "DOB", "Lead_Creation_Date", "City_Code",
    ↪ "Employer_Code", "Customer_Existing_Primary_Bank_Code", "Source"])

#Updating categorical columns:
categorical_columns = ["Gender", "City_Category", "Employer_Category1",
    ↪ "Employer_Category2", "Primary_Bank_Type", "Contacted", "Source_Category"]

```

Feature Engineering: I will now create the Total_Loan_Amount column which is the loan amount for the entire loan period including interest rate. This column will be used in A/B testing.

```

[63]: df["Total_Loan_Amount"] = df["Loan_Amount"] * (1 + (df["Interest_Rate"]/100) *
    ↪ df["Loan_Period"])
df["Total_Loan_Amount"].describe()

#Updating the numerical columns
numerical_columns = ["Monthly_Income", "Existing_EMI", "Loan_Amount",
    ↪ "Loan_Period", "Interest_Rate", "EMI", "Age", "Lead_Age",
    ↪ "Total_Loan_Amount"]

```

Visualizations:

```

[64]: #Grid layout:
n_cols = 3
n_rows = 3

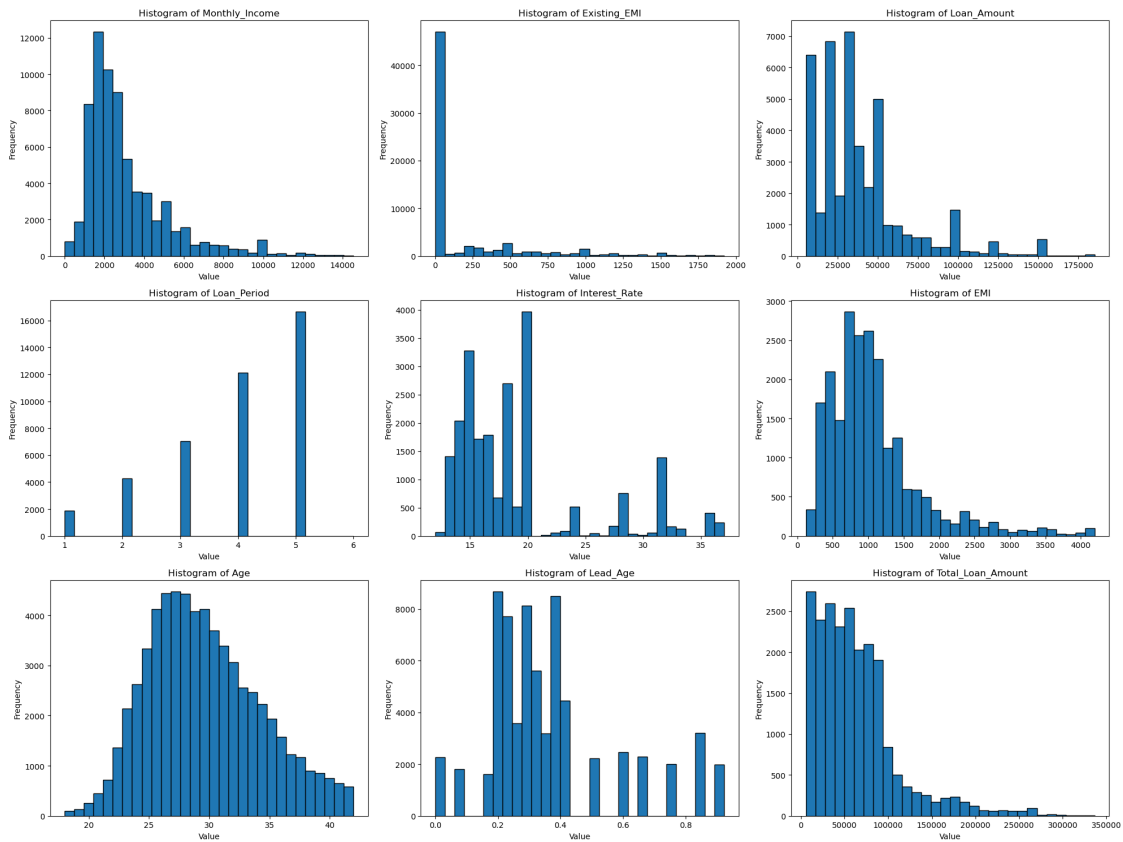
plt.figure(figsize=(20, 15))

for i, col in enumerate(numerical_columns, 1):
    plt.subplot(n_rows, n_cols, i)
    plt.hist(df[col], bins=30, edgecolor="black")
    plt.title(f"Histogram of {col}")

```

```
plt.xlabel("Value")
plt.ylabel("Frequency")
```

```
plt.tight_layout()
plt.show()
```



```
[65]: n_cols = 3
n_rows = 3

plt.figure(figsize=(20, 15))

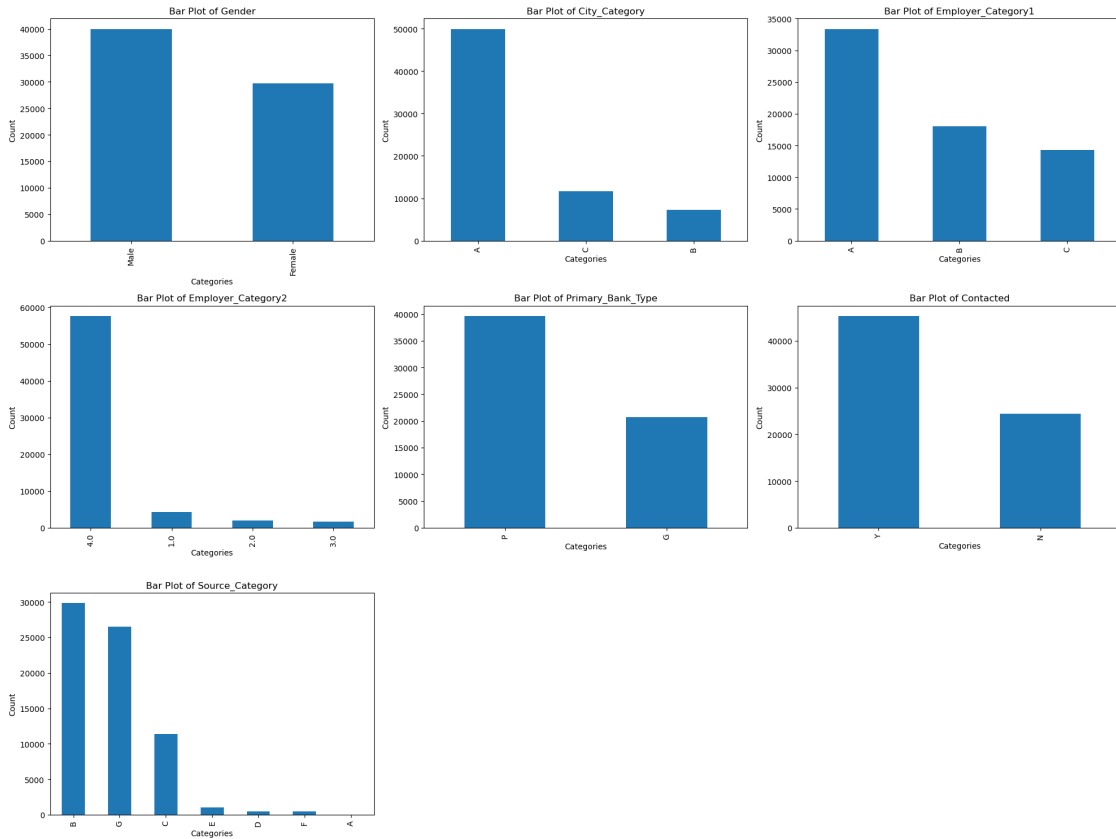
for i, col in enumerate(categorical_columns, 1):
    plt.subplot(n_rows, n_cols, i)

    value_counts = df[col].value_counts()

    value_counts.plot(kind="bar")
    plt.title(f"Bar Plot of {col}")
    plt.xlabel("Categories")
    plt.ylabel("Count")
```



```
plt.tight_layout()
plt.show()
```



2 Statistical Tests:

For the initial A/B tests I will explore the approval rate amongst the lead source categories, to determine which categories are the most effective.

```
[66]: approval_rates = df.groupby("Source_Category")["Approved"].mean() * 100
```

We can see that there are 7 distinct source categories. B and G are the most commonly occurring categories by a substantial amount. However, category E which is the fourth most common source category has the second highest approval rating. Therefore it is worthwhile to explore whether the approval rating of category E is significantly greater than that of category G.

First however, I will test the significance of the difference between the approval ratings of categories B and G, the two most commonly occurring categories.

2.0.1 Research Hypothesis:

Null Hypothesis (H0): There is no statistically significant difference in the loan approval rate between the “G” and “B” lead source categories.

Alternative Hypothesis (H1): There is a statistically significant difference in the loan approval rate between the “G” and “B” lead source categories.

2.0.2 A/B Testing Using the Chi-Square Test

For this series of A/B tests I will use the Chi-Square test. Since we are dealing with categorical data and there are many data points, the Chi-Square test is the most appropriate.

```
[67]: #Setting the significance level at a standard value of 0.05
alpha = 0.05

#Filtering dataset
df_filtered = df[df["Source_Category"].isin(["G", "B"])]

#Calculating approval rates
approval_rates = df_filtered.groupby("Source_Category")["Approved"].mean() * 100
print(approval_rates)

#Creating contingency table needed for the Chi-Square test
contingency_table = pd.crosstab(df_filtered["Source_Category"],
    ↪df_filtered["Approved"])

chi2, p_value, dof, expected = chi2_contingency(contingency_table)

print(f"\nChi-Square Statistic: {chi2}")
print(f"P-value: {p_value}")

if p_value < alpha:
    print("\nReject the null hypothesis: There is a significant difference in
    ↪approval rates between the B and G source categories.")
else:
    print("\nFail to reject the null hypothesis: There is no significant
    ↪difference in approval between the B and G source categories.")
```

```
Source_Category
B      1.660405
G      1.372653
Name: Approved, dtype: float64
```

```
Chi-Square Statistic: 7.547966987422923
P-value: 0.006007781759069996
```

Reject the null hypothesis: There is a significant difference in approval rates between the B and G source categories.

The Chi-Square test yielded a p-value of approximately 0.006. This is less than the set significance level of 0.05, which means we reject the null hypothesis and conclude that there is indeed a significant difference in approval rates in leads between the B and G source categories.

Now let's see if source category E is also significantly greater than G.

Null Hypothesis (H0): There is no statistically significant difference in the loan approval rate between the "G" and "E" lead source categories.

Alternative Hypothesis (H1): There is a statistically significant difference in the loan approval rate between the "G" and "E" lead source categories.

```
[68]: df_filtered = df[df["Source_Category"].isin(["G", "E"])]

approval_rates = df_filtered.groupby("Source_Category")["Approved"].mean() * 100
print(approval_rates)

contingency_table = pd.crosstab(df_filtered["Source_Category"],
                                df_filtered["Approved"])

chi2, p_value, dof, expected = chi2_contingency(contingency_table)

print(f"\nChi-Square Statistic: {chi2}")
print(f"P-value: {p_value}")

if p_value < alpha:
    print("\nReject the null hypothesis: There is a significant difference in
    approval rates between the E and G source categories.")
else:
    print("\nFail to reject the null hypothesis: There is no significant
    difference in approval between the E and G source categories.")
```

```
Source_Category
E    1.428571
G    1.372653
Name: Approved, dtype: float64
```

```
Chi-Square Statistic: 0.0003064824842084836
P-value: 0.9860324347908929
```

Fail to reject the null hypothesis: There is no significant difference in approval between the E and G source categories.

We failed to reject the null hypothesis, therefore there is not enough conclusive evidence to suggest that the E source category has a significantly different approval rate than the G category.

Looking at the bar plot we can see that the C source category also has a lower approval rate than E, while having many more occurrences than E. Therefore it may be worthwhile to investigate whether or not E's approval rating is significantly different.

Null Hypothesis (H0): There is no statistically significant difference in the loan approval rate

between the “C” and “E” lead source categories.

Alternative Hypothesis (H1): There is a statistically significant difference in the loan approval rate between the “C” and “E” lead source categories.

```
[69]: #Filtering dataset to only include G and B source categories.
df_filtered = df[df["Source_Category"].isin(["C", "E"])]

#Calculating approval rates for both categories:
approval_rates = df_filtered.groupby("Source_Category")["Approved"].mean() * 100
print(approval_rates)

#Creating contingency table needed for the Chi-Square test
contingency_table = pd.crosstab(df_filtered["Source_Category"],
    ↪df_filtered["Approved"])

print("Contingency Table:")
print(contingency_table)

chi2, p_value, dof, expected = chi2_contingency(contingency_table)

print(f"\nChi-Square Statistic: {chi2}")
print(f"P-value: {p_value}")

if p_value < alpha:
    print("\nReject the null hypothesis: There is a significant difference in
    ↪approval rates between the E and C source categories.")
else:
    print("\nFail to reject the null hypothesis: There is no significant
    ↪difference in approval between the E and C source categories.")
```

```
Source_Category
C    1.230877
E    1.428571
Name: Approved, dtype: float64
Contingency Table:
Approved          0    1
Source_Category
C             11234  140
E             1035   15
```

```
Chi-Square Statistic: 0.16558336652502456
P-value: 0.6840672153648855
```

Fail to reject the null hypothesis: There is no significant difference in approval between the E and C source categories.

There is not enough conclusive evidence to suggest that the E category is significantly larger than the G or C categories. Therefore, the B and G categories are still the most effective lead generation

categories.

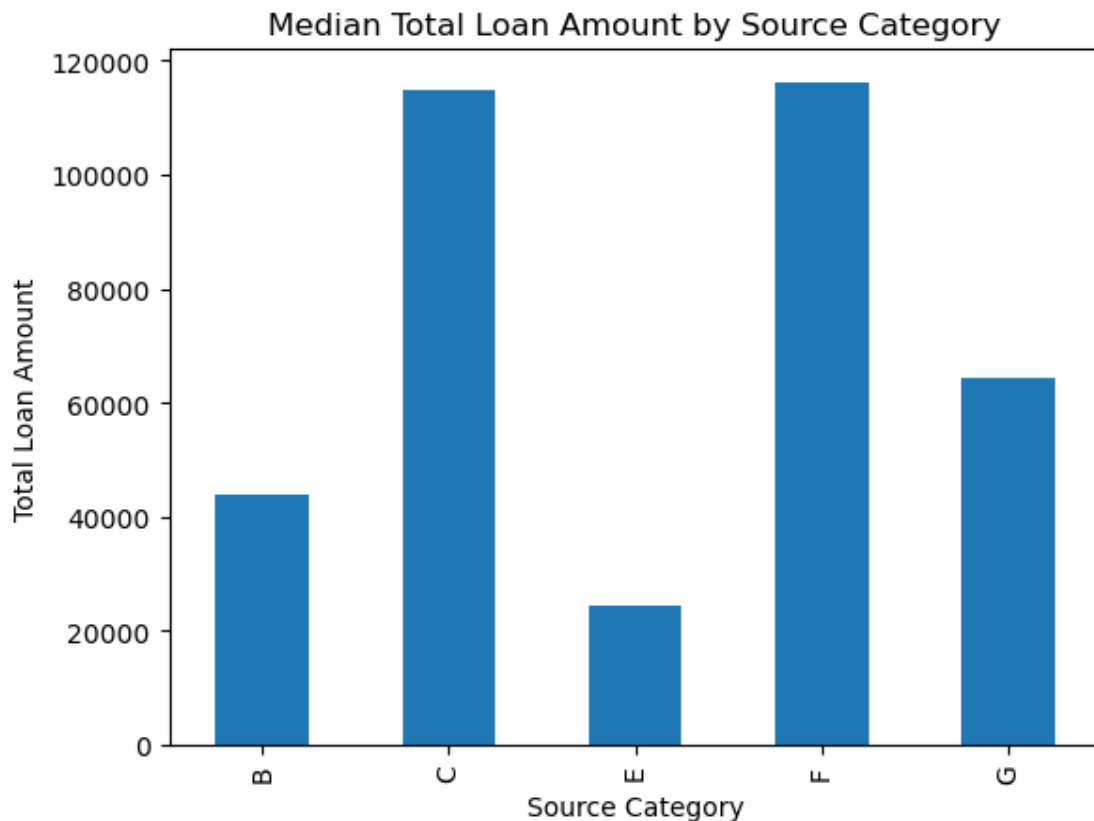
2.0.3 A/B Testing Total Loan Amount Using the Mann-Whitney U test:

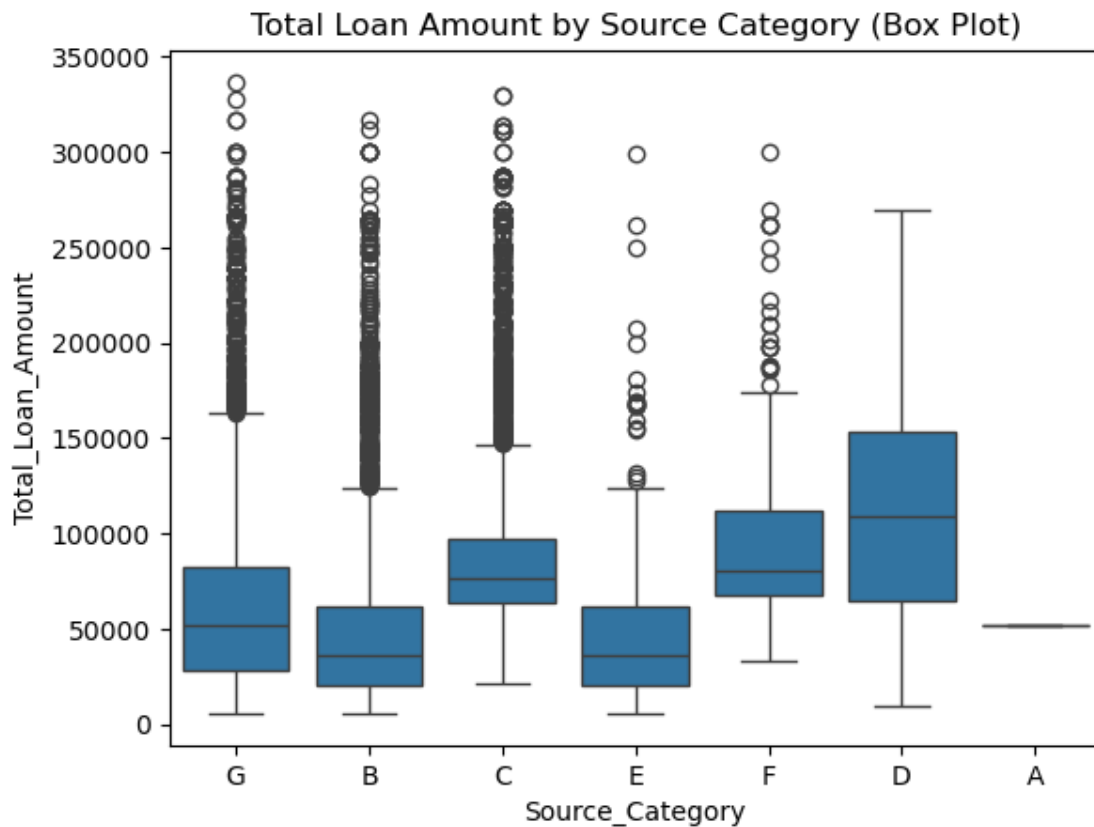
Next I will perform an A/B test to determine if the G source category has a significantly greater median Total Loan Amount than the B source category. Since I am now comparing continuous data (Total Loan Amount), I have the option of choosing the T-test or the Mann-Whitney U test. This decision depends on the distribution of the data. If the data is normally distributed, the T-test will be applicable, if not, then the Mann-Whitney U test will be more appropriate.

```
[70]: approved_df = df[df["Approved"] == 1]
total_loans = approved_df.groupby("Source_Category")["Total_Loan_Amount"].
    ↪median()

# Plot the approval rates
total_loans.plot(kind="bar")
plt.title("Median Total Loan Amount by Source Category")
plt.ylabel("Total Loan Amount")
plt.xlabel("Source Category")
plt.show()

sns.boxplot(x='Source_Category', y='Total_Loan_Amount', data=df)
plt.title('Total Loan Amount by Source Category (Box Plot)')
plt.show()
```





```
[71]: cat_G = df[df["Source_Category"].isin(["G"])]
cat_B = df[df["Source_Category"].isin(["B"])]

cat_G = cat_G.dropna(subset=["Total_Loan_Amount"])
cat_B = cat_B.dropna(subset=["Total_Loan_Amount"])

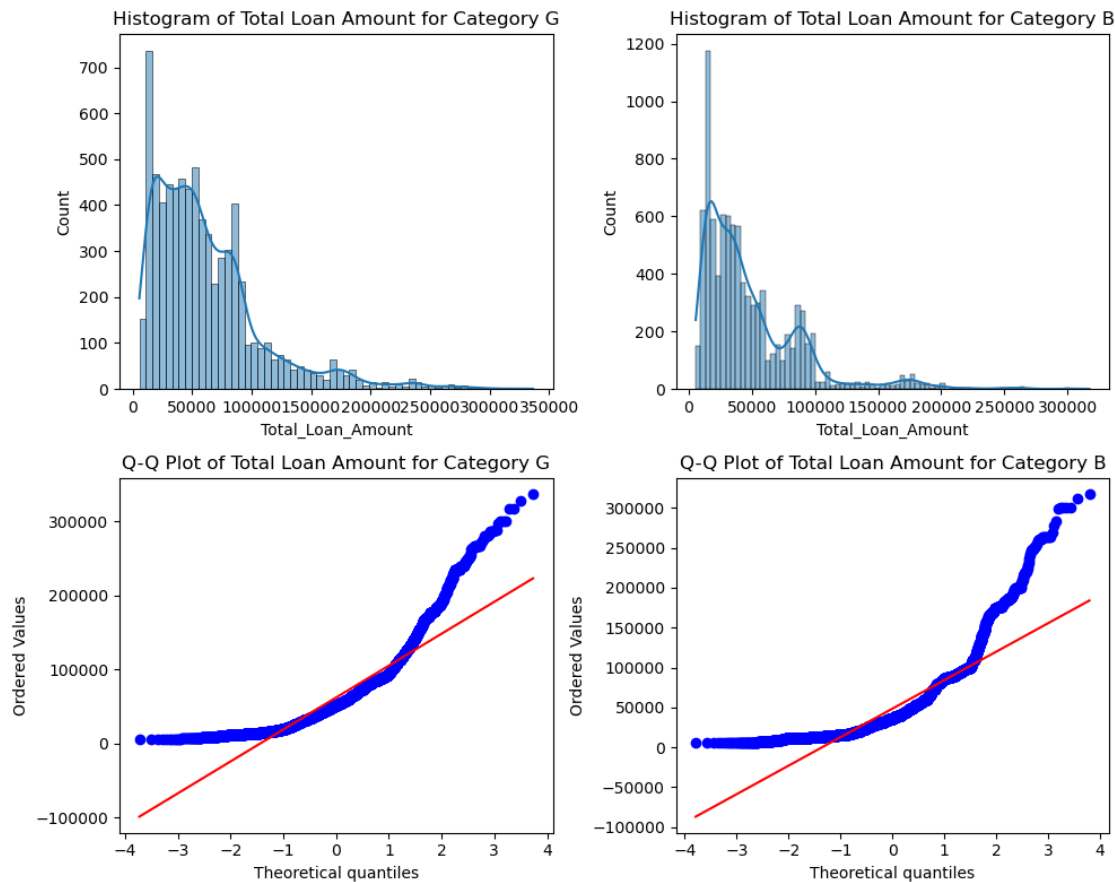
fig, axes = plt.subplots(2, 2, figsize=(10, 8))

#Histograms of Total Loan Amount
sns.histplot(cat_G["Total_Loan_Amount"], kde=True, ax=axes[0,0])
axes[0,0].set_title("Histogram of Total Loan Amount for Category G")
sns.histplot(cat_B["Total_Loan_Amount"], kde=True, ax=axes[0,1])
axes[0,1].set_title("Histogram of Total Loan Amount for Category B")

#Q-Q Plots of Total loan amount
stats.probplot(cat_G["Total_Loan_Amount"], dist="norm", plot=axes[1,0])
axes[1,0].set_title("Q-Q Plot of Total Loan Amount for Category G")
stats.probplot(cat_B["Total_Loan_Amount"], dist="norm", plot=axes[1,1])
```

```
axes[1,1].set_title("Q-Q Plot of Total Loan Amount for Category B")

plt.tight_layout()
plt.show()
```



From the above plots we can see the data is heavily skewed and does not follow a normal distribution. This is especially apparent in the Q-Q plots. Therefore, I will perform the Mann-Whitney U test.

```
[72]: u_stat, p_value = mannwhitneyu(cat_G["Total_Loan_Amount"],
    ↪ cat_B["Total_Loan_Amount"], alternative="two-sided")

print(f"P-value = {p_value}")

if p_value < alpha:
    print("Reject the null hypothesis, there is a significant difference in the
    ↪ Total Loan Amount between the two groups.")
else:
    print("Fail to reject the null hypothesis, there is no significant
    ↪ difference in the Total Loan Amount between the two groups.")
```

P-value = 1.346446553936553e-106

Reject the null hypothesis, there is a significant difference in the Total Loan Amount between the two groups.

Since the P-value is significantly below the significance level of 0.05, we reject the null hypothesis. Therefore the G category has a significantly greater Total Loan Amount.

3 Logistic Regression:

I will now build a predictive model using logistic regression to predict approval of a bank loan application. Before doing so I must encode the categorical columns into binary columns using One Hot Encoding, as previously mentioned.

The null values will also be imputed using the KNN algorithm. The K-nearest neighbours algorithm will impute null values based on the mean of the 5 closest neighbours of the row in question. Before doing so, I must scale the numerical values as the vast difference in ranges between variables will result in a poor quality Euclidean distance calculation in the KNNImputer method.

```
[73]: numerical_columns = ["Monthly_Income", "Existing_EMI", "Loan_Amount",  
    ↪ "Loan_Period", "Interest_Rate", "EMI", "Total_Loan_Amount", "Age",  
    ↪ "Lead_Age"]  
categorical_columns = ["Gender", "City_Category", "Employer_Category1",  
    ↪ "Employer_Category2", "Primary_Bank_Type", "Contacted", "Source_Category"]  
df_log = df  
  
#Using OneHotEncoder to convert categorical columns to numerical  
encoder = OneHotEncoder(sparse_output=False, handle_unknown='error')  
encoded_columns = encoder.fit_transform(df_log[categorical_columns])  
  
#Add encoded columns back to the DataFrame  
encoded_df = pd.DataFrame(encoded_columns, columns=encoder.  
    ↪get_feature_names_out(categorical_columns))  
df_log = pd.concat([df.drop(columns=categorical_columns), encoded_df], axis=1)  
  
print(df_log.dtypes)  
  
#KNN Imputing:  
  
#Scaling numerical columns  
scaler = StandardScaler()  
scaled_columns = scaler.fit_transform(df_log[numerical_columns])  
  
scaled_df = pd.DataFrame(scaled_columns, columns=scaler.  
    ↪get_feature_names_out(numerical_columns))  
df_log = pd.concat([df_log.drop(columns=numerical_columns), scaled_df], axis=1)  
  
imputer = KNNImputer(n_neighbors=5)
```



```

imputed_data = imputer.fit_transform(df_log)

# Convert the result back to a DataFrame
df_log = pd.DataFrame(imputed_data, columns=df_log.columns)
print("DataFrame After Imputation:")
print(df_log)

```

```

Monthly_Income      float64
Existing_EMI         float64
Loan_Amount          float64
Loan_Period          float64
Interest_Rate        float64
EMI                  float64
Approved             int64
Lead_Age             float64
Age                  float64
Total_Loan_Amount    float64
Gender_Female        float64
Gender_Male          float64
City_Category_A      float64
City_Category_B      float64
City_Category_C      float64
City_Category_nan    float64
Employer_Category1_A float64
Employer_Category1_B float64
Employer_Category1_C float64
Employer_Category1_nan float64
Employer_Category2_1.0 float64
Employer_Category2_2.0 float64
Employer_Category2_3.0 float64
Employer_Category2_4.0 float64
Employer_Category2_nan float64
Primary_Bank_Type_G  float64
Primary_Bank_Type_P  float64
Primary_Bank_Type_nan float64
Contacted_N          float64
Contacted_Y          float64
Source_Category_A    float64
Source_Category_B    float64
Source_Category_C    float64
Source_Category_D    float64
Source_Category_E    float64
Source_Category_F    float64
Source_Category_G    float64
dtype: object

```

DataFrame After Imputation:

	Approved	Gender_Female	Gender_Male	City_Category_A	City_Category_B	\
0	0.0	1.0	0.0	1.0	0.0	

1	0.0	0.0	1.0	1.0	0.0
2	0.0	0.0	1.0	0.0	0.0
3	0.0	0.0	1.0	0.0	0.0
4	0.0	0.0	1.0	1.0	0.0
...
69708	0.0	1.0	0.0	1.0	0.0
69709	0.0	1.0	0.0	0.0	0.0
69710	0.0	1.0	0.0	0.0	1.0
69711	0.0	0.0	1.0	1.0	0.0
69712	0.0	0.0	1.0	1.0	0.0

	City_Category_C	City_Category_nan	Employer_Category1_A \
0	0.0	0.0	1.0
1	0.0	0.0	0.0
2	1.0	0.0	0.0
3	1.0	0.0	1.0
4	0.0	0.0	1.0
...
69708	0.0	0.0	1.0
69709	1.0	0.0	1.0
69710	0.0	0.0	0.0
69711	0.0	0.0	0.0
69712	0.0	0.0	1.0

	Employer_Category1_B	Employer_Category1_C	...	Source_Category_G \
0	0.0	0.0	...	1.0
1	0.0	1.0	...	1.0
2	0.0	1.0	...	0.0
3	0.0	0.0	...	0.0
4	0.0	0.0	...	0.0
...
69708	0.0	0.0	...	1.0
69709	0.0	0.0	...	1.0
69710	0.0	1.0	...	1.0
69711	0.0	1.0	...	1.0
69712	0.0	0.0	...	1.0

	Monthly_Income	Existing_EMI	Loan_Amount	Loan_Period	Interest_Rate \
0	-0.487853	-0.513225	0.114117	0.093681	0.356495
1	0.218947	-0.513225	-0.650051	-1.619415	-1.019936
2	-0.370053	-0.513225	0.210499	0.093681	-0.532507
3	0.218947	-0.513225	1.828332	0.950229	-0.019424
4	3.281748	0.530382	0.382609	-1.619415	-0.599208
...
69708	0.878627	-0.513225	-0.168143	-1.619415	-0.815387
69709	1.957723	3.269851	-0.725780	-0.934176	-0.430574
69710	-0.676333	-0.513225	-0.512363	0.093681	2.785431
69711	3.231330	3.050694	1.415268	0.950229	-0.849592

69712	0.562923	-0.513225	1.036626	0.093681	-0.893375
-------	----------	-----------	----------	----------	-----------

	EMI	Total_Loan_Amount	Age	Lead_Age
0	-0.150087	0.052283	1.669508	0.181767
1	-0.180387	-0.805973	0.176668	1.490864
2	-0.200187	-0.118003	0.979377	0.128874
3	2.018915	2.505497	-0.370393	-0.532287
4	0.766414	-0.208451	0.439938	0.115651
...
69708	0.852514	-0.414438	0.808164	-0.836420
69709	-0.008787	-0.179195	0.906201	-0.836420
69710	-0.195387	-0.113248	-0.996025	-0.836420
69711	1.024114	1.375698	1.898769	-0.836420
69712	1.217615	0.933497	-0.566819	-0.836420

[69713 rows x 37 columns]

Before proceeding to the logistic regression, I will analyze the VIF (Variance Inflation Factor) scores of the predictor variables. VIF scores are used to detect multicollinearity among the independent variables, if any of the VIF scores are too high (above 10 is the standard threshold) we will know that some of the variables have high multicollinearity with other variables. Highly multicollinear variables risk overfitting the model, reducing the model's ability to generalize to new data.

```
[74]: #VIFS analysis:
predictor_variables = df_log.drop(columns=['Approved'])

vifs = pd.DataFrame()
vifs['Feature'] = predictor_variables.columns
vifs['VIF'] = [variance_inflation_factor(predictor_variables.values, i) for i
               in range(predictor_variables.shape[1])]

print("VIF before regression:")
print(vifs)
```

VIF before regression:

	Feature	VIF
0	Gender_Female	5.597318e+03
1	Gender_Male	1.239562e+05
2	City_Category_A	4.289051e+04
3	City_Category_B	3.102692e+04
4	City_Category_C	8.636364e+04
5	City_Category_nan	1.300771e+10
6	Employer_Category1_A	3.164695e+04
7	Employer_Category1_B	2.606334e+04
8	Employer_Category1_C	1.376348e+05
9	Employer_Category1_nan	7.217196e+06
10	Employer_Category2_1.0	1.696727e+05
11	Employer_Category2_2.0	9.477563e+07

```

12 Employer_Category2_3.0 1.190566e+09
13 Employer_Category2_4.0 2.004711e+03
14 Employer_Category2_nan 5.719249e+06
15 Primary_Bank_Type_G 1.082611e+05
16 Primary_Bank_Type_P 2.390920e+05
17 Primary_Bank_Type_nan 1.323379e+06
18 Contacted_N 3.671566e+04
19 Contacted_Y 2.383195e+04
20 Source_Category_A 2.912124e+12
21 Source_Category_B 1.248028e+04
22 Source_Category_C 5.316565e+05
23 Source_Category_D 6.347908e+09
24 Source_Category_E 3.271790e+06
25 Source_Category_F 1.949933e+07
26 Source_Category_G 4.427334e+05
27 Monthly_Income 2.267830e+00
28 Existing_EMI 1.259897e+00
29 Loan_Amount 1.250151e+01
30 Loan_Period 2.878793e+00
31 Interest_Rate 1.782132e+00
32 EMI 1.308810e+01
33 Total_Loan_Amount 1.850744e+01
34 Age 1.133724e+00
35 Lead_Age 1.064465e+00

```

As we can see, many of VIF scores are extremely high, indicating that many of the variables have multicollinearity issues. I will create a while loop that will iteratively remove the column with the highest VIF value and recalculate the VIFs until there are no columns left with VIF scores over 10.

```

[75]: while(vifs['VIF'].max() > 10):
    max_VIF = vifs['VIF'].max()
    dropped_column = vifs.loc[vifs['VIF'] == max_VIF, 'Feature'].iloc[0]

    print(f"Dropping {dropped_column}, VIF: {max_VIF}")
    df_log = df_log.drop(columns=[dropped_column])
    predictor_variables = predictor_variables.drop(columns=[dropped_column])

    vifs = pd.DataFrame()
    vifs['Feature'] = predictor_variables.columns
    vifs['VIF'] = [variance_inflation_factor(predictor_variables.values, i) for
    ↪ i in range(predictor_variables.shape[1])]

```

```

Dropping Source_Category_A, VIF: 2912123910359.196
Dropping City_Category_nan, VIF: 65332525221.70637
Dropping Employer_Category2_2.0, VIF: 516280162.75854945
Dropping Primary_Bank_Type_nan, VIF: 654160586.0618285
Dropping Employer_Category1_B, VIF: 412139222.2230963

```

```
c:\Users\dimdi\anaconda3\Lib\site-
```

```
packages\statsmodels\stats\outliers_influence.py:197: RuntimeWarning: divide by
zero encountered in scalar divide
```

```
vif = 1. / (1. - r_squared_i)
```

```
Dropping Gender_Female, VIF: inf
```

```
Dropping Contacted_Y, VIF: 15197.369998117312
```

```
Dropping City_Category_A, VIF: 69.47832812300553
```

```
Dropping Employer_Category2_4.0, VIF: 31.67279008450572
```

```
Dropping Total_Loan_Amount, VIF: 18.51951991167427
```

```
Dropping Employer_Category1_nan, VIF: 16.142612890802386
```

```
Dropping Source_Category_B, VIF: 12.9248116137109
```

```
[76]: print("VIF before regression:")
      print(vifs)
```

```
VIF before regression:
```

	Feature	VIF
0	Gender_Male	5.350031
1	City_Category_B	1.187167
2	City_Category_C	1.359705
3	Employer_Category1_A	2.772578
4	Employer_Category1_C	1.775623
5	Employer_Category2_1.0	1.135784
6	Employer_Category2_3.0	1.047284
7	Employer_Category2_nan	1.574947
8	Primary_Bank_Type_G	3.166814
9	Primary_Bank_Type_P	4.872251
10	Contacted_N	3.679494
11	Source_Category_C	1.803486
12	Source_Category_D	1.063761
13	Source_Category_E	1.035206
14	Source_Category_F	1.032498
15	Source_Category_G	2.228823
16	Monthly_Income	2.259405
17	Existing_EMI	1.257358
18	Loan_Amount	9.436102
19	Loan_Period	1.871734
20	Interest_Rate	1.733664
21	EMI	8.131278
22	Age	1.131600
23	Lead_Age	1.055468

Now that the multicollinearity issue is dealt with, the data is ready for logistic regression. I will split the dataframe into separate testing and training data sets.

```
[77]: X = predictor_variables
      y = df_log["Approved"]
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1,
↳stratify=y)
```

Now I can fit the logistic regression to the training data and evaluate it using the testing data.

```
[78]: model = LogisticRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print(classification_report(y_test, y_pred))
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

roc_auc_scores = cross_val_score(model, X, y, cv=5, scoring='roc_auc')
print(f"Mean ROC-AUC: {roc_auc_scores.mean()}")
print(f"Standard Deviation of ROC-AUC: {roc_auc_scores.std()}")

coef = model.coef_[0]
predictors = X.columns
plt.barh(predictors, coef)
plt.title('Predictor Coefficients')
plt.show()
```

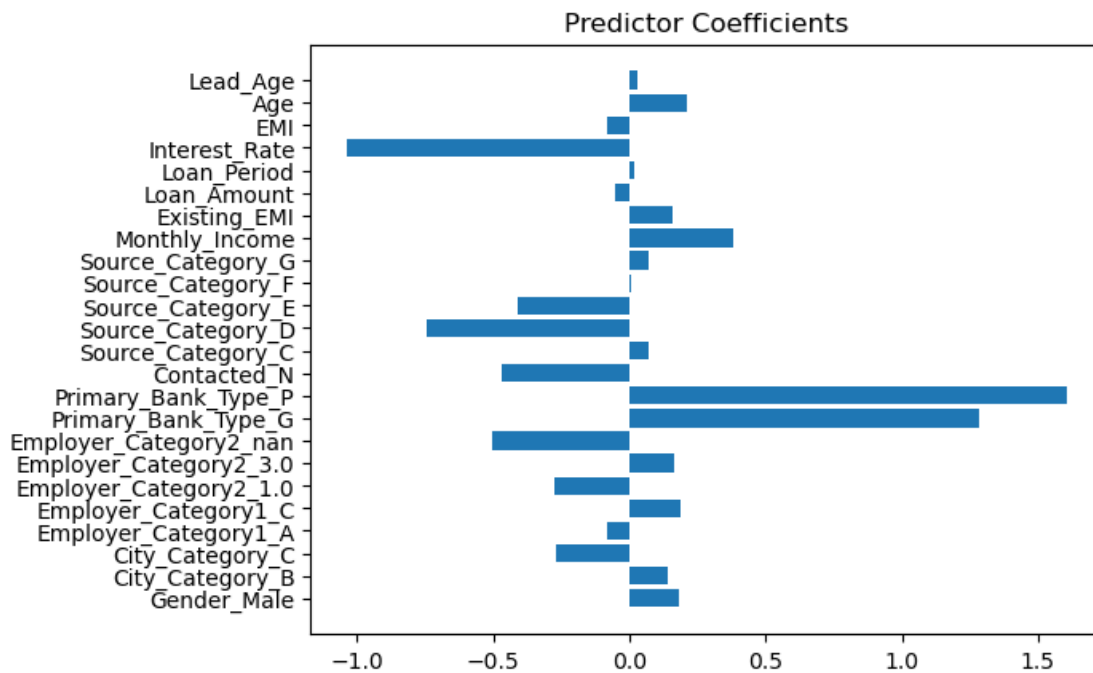
```
c:\Users\dimdi\anaconda3\Lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\dimdi\anaconda3\Lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\dimdi\anaconda3\Lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	6870
1.0	0.00	0.00	0.00	102
accuracy			0.99	6972
macro avg	0.49	0.50	0.50	6972
weighted avg	0.97	0.99	0.98	6972

Accuracy: 0.9853700516351118

Mean ROC-AUC: 0.8109390287077209

Standard Deviation of ROC-AUC: 0.014821339432588749



We can see that the predictive model worked well, with approximately a 99% success rate.

The mean ROC-AUC is 0.81, meaning that the model has good predictive power. A ROC-AUC score of 0.5 would indicate that the model simply makes random guesses, a ROC-AUC score of 1.0 indicates perfect predictive power. Therefore a ROC-AUC score of 0.81 indicates that it can effectively distinguish the outcome.

We can also see that the most important predictors (highest absolute magnitude) are the two primary bank types and the interest rate. The strong positive coefficient of the primary bank types would suggest that customers from those bank types are much more likely to get approved. The high negative coefficient for interest rate would suggest that higher interest rate loans are less likely to be approved.

Monthly income was the third most important variable with a positive coefficient in predicting approval, meaning that higher montly incomes are more likely to get their loan approved, as would be expected.